One Initialization to Rule Them All

Stefan Wijnja, Ellis Wierstra, Thomas Hamburger

January 27, 2020

Neural networks are very cool

- Images of cool applications.

But they can be very demanding

- Numbers for huge networks and how long it takes to train them.

One possible solution: pruning

- Explain pruning, iterative pruning, discuss known results so far.

Pruning can lead to Winning Tickets

- Explain what winning tickets are.

Question

Does weight initialization matter looking for Winning Tickets?

Method

Model

- Fully Connected.
- Two hidden layers: 300 & 100 neurons $\rightarrow$ 266k weights.

Training

Pruning

- Used LeNet on MNIST and tried a bunch of different inits. Nice summary on framework(s) used, how many lines of code written, how many epochs trained, etc.>

Results

- Nice graphs of pruning rate vs. test_acc.

Further Research

- What does this research lead to?

Questions?